

---

# Object Oriented Languages

---

CHAPTERS 7, 8, AND 10 OF PROGRAMMING LANGUAGES  
PRAGMATICS

---

# Three fundamental concepts

---

- **Encapsulation:**

- Group data and subroutines in one place, hide irrelevant details from users of the abstraction

- **Inheritance:**

- New abstractions defined as refinements or extensions of existing abstractions

- **Dynamic Method Binding:**

- Allow a new version of an abstraction to display newly refined behavior, even when used in a context that expects an earlier version

# Outline

---

- **Classes and objects**
- **Methods**
  - Inheritance, polymorphism
  - Static methods and fields
- **Implementation: compilation, allocation**
- **Types**
- **Memory and type safety**
- **Memory Management**

# Classes

---

- A **class** is a blueprint for creating objects
  - Contains all the information the programmer needs to use the abstraction
  - Contains all the information the compiler needs to generate code

```
class Rectangle {  
    public double height, width;  
    public double area() {  
        return height * width;  
    }  
}
```

- This is Java code, the equivalent C++ code is very similar
- Class members: **methods** and **fields**

# Objects

---

- The central concept of object oriented programming
- In C++ and Java, objects are **instances** of classes, created through **new**
  - For example, when the expression **new** Rectangle() is evaluated, a new object of class Rectangle is created and initialized
  - “instance” = “object”
  - “class X is instantiated” = “an instance of X is created”

# Creation of objects

---

- During the evaluation of `x = new Rectangle()`
  - A new instance (object) of class Rectangle is created on the heap
  - A reference (pointer) to this new instance is produced
    - This is the result of evaluating the new expression
  - The appropriate constructor of the class is called to initialize the new object
  - `x` is assigned this reference (pointer) value
    - the value may be the address of the first byte of the objects memory
    - or the value may be some internal handle to the actual object (i.e. an index in some internal table which contains the address of the first byte)

# Destruction of objects

---

- C++: each new must have a corresponding delete
  - `x = new Rectangle; ... delete x;`
- Java: dead objects are reclaimed automatically by a garbage collector (GC)
  - `x = new Rectangle();` // After you stop using the object, the GC (hopefully) figures out it is dead
- C++ destructors: called when the programmer manually destroys the object with delete
  - `class Rectangle { ... ~Rectangle() { ... } }`
- Java finalizers: called when the object is collected
  - `class Rectangle { ... void finalize() { ... } }`

# Members: fields and methods

---

- Two separate kinds:
  - **Instance** members
    - Each **instance** of the class has a separate copy of this member
  - **Static** members

```
Rectangle a, b, c;  
a = new Rectangle();  
b = new Rectangle();  
a.height = 1.0; a.width = 3.6;  
b.height = 2.2; b.width = 5.0;  
c = a;
```



# Members: fields and methods (C++)

---

- C++: `x->f` is shorthand for `(*x).f`
  - `x` evaluates to a pointer value that points to the object
  - `*x` evaluates to the actual object
  - `(*x).f` evaluates to the field `f` of the object

```
Rectangle *a, *b, *c;  
a = new Rectangle();  
b = new Rectangle();  
a->height = 1.0; a->width = 3.6;  
b->height = 2.2; b->width = 5.0;  
c = a;
```

# Outline

---

- Classes and objects
- **Methods**
  - Inheritance, polymorphism
  - Static methods and fields
- Implementation: compilation, allocation
- Types
- Memory and type safety
- Memory Management

# Instance methods

---

- An instance method operates on objects
  - Is invoked on the object

```
double area() {return height * width; }
```

- In reality, this is shorthand for

```
double area(Rectangle this) {  
    return this.height * this.width; } //Java
```

```
double area(Rectangle* this) {  
    return this->height * this->width; } //C++
```

- There is an implicit formal parameter this: a reference to the object on which the method was invoked
  - Calls like `x.area()` and `x->area()` are, essentially, calls to `area(x)`

# Constructors

---

- Constructors are used to set up the initial state of new objects

```
public Rectangle(double height, double width) {  
    this.height = height; this.width = width;  
}
```

- `x = new Rectangle(1.1, 2.3);`
  - A new object is created, with default value 0.0 in Java and undefined values in C++
  - The constructor is invoked on this object; the fields are initialized with 1.1 and 2.3
  - A reference of the object is assigned to `x`

# Outline

---

- Classes and objects
- Methods
  - **Inheritance, polymorphism**
  - Static methods and fields
- Implementation: compilation, allocation
- Types
- Memory and type safety
- Memory Management

# Inheritance

---

- class **B** extends **A** { ... }
  - Single inheritance: only one superclass (Java)
- class **B** : public **A1**, **A2**, **A3** { ... }
  - Multiple inheritance: several superclasses (C++)
- Every member of **A** is inherited by **B**
  - If a field *f* is defined in **A**, every object of class **B** has an *f* field
  - If a method is defined in **A**, this method can be invoked on an object of class **B**
- **B** may declare new members

# Example

---

```
class Rectangle {  
    private double height, width;  
    public double getHeight() { return height; }  
    public double get Width() { return width; }  
    public double area() { ... }  
}  
  
class SwissRectangle extends Rectangle {  
    private int hole_size;  
    public SwissRectangle(double h, double w, int hs) { ... }  
    public void shrinkHold() { hole_size--; }  
    public double area() { ... }  
}
```

# Constructors and inheritance

---

- Constructors are not inherited
- A constructor in a subclass **B** must invoke a constructor in the superclass **A**
  - This is maybe an oversimplification
- The constructor of superclass **A** initialized the part of the “object state” that is declared in **A**
  - Sets up values for fields declared in **A** and inherited by the subclasses

```
class SwissRectangle extends Rectangle {  
    private int hole_size;  
    public SwissRectangle(double h, double w, int hs) {  
        super(h, w);  
        hole_size = hs;  
    }  
}
```



# Inheritance of methods

---

- If a subclass declares a method with the same name but a different signature, we have **overloading**
  - Method signature: Some combination of
    - name
    - number and type of arguments
    - return type
  - **Either** method can be invoked on an instance of the subclass
- If a subclass declares a method with the same signature, we have **overriding**
  - **Only** the new method applies to instances of the subclass

# Polymorphism of references

---

- Reference variables for **A** objects also may point to **B** objects
  - **A** **x** = new **B**() in Java
  - **A**\* **x** = new **B**() in C++
- Simplistic view: the type of **x** is pointer (reference) to instance of **A**
- Correct view: pointer (reference) to instance of **A** or instances of any transitive subclass of **A**
  - If **C** is a subclass of **B**, variable **x** can also point to instance of **C**
    - **C** is a transitive subclass of **A**
  - **Poly** (many) **morph** (form) **ism** (state or condition)

# Method invocation – compile time

---

- What happens when we have a method invocation of the form `x.m(a, b)`?
- Two very different things are done
  - At **compile time**, by the Java compiler (javac)
  - At **run time**, by the Java Virtual Machine
- At **compile time**, a target method is associated with the invocation expression
  - Terms: **compile-time target**, **static target**
  - The **static target** is **based on the declared type of x**

# Method invocation – compile time

---

```
class A {void m(int p, int q) { ... } ... }  
class B extends A { void m(int r, int s) { ... } ... }  
A x;  
x = new B();  
x.m(1, 2);
```

- Since **x** has declared type **A**, the compile-time target is method **m** in class **A**
- Javac encodes this in the bytecode
  - virtualinvoke **x**,<**A**: void m(**int**,**int**)>

# Method invocation – run time

---

- The Java Virtual Machine loads the bytecode and starts executing it
- When the JVM tries to execute the instruction `virtualinvoke x,<A: void m(int,int)>`
  - JVM looks at the class **Z** of the object referenced by **x**
    - i.e. what **x** actually is at run time
  - JVM searches **Z** for a method with signature `m(int,int)` and return type `void`
    - Searches for a match to the signature
  - If **Z** does not have it, the JVM goes to **Z**'s superclass to find a match, and continues to search upwards through the class hierarchy until a match is found?
    - Is the JVM guaranteed to find a match?

# Method invocation – run time

---

- The run time (**dynamic**) target: The “lowest” method that matches the signature of the static target method
  - “Lowest” with respect to the inheritance chain from **Z** to **java.lang.Object**
- Once the JVM determines the run time target method, it invokes it on the object that is referenced by **x**
- Key term:
  - **Virtual dispatch/dynamic dispatch**: The process of selecting which implementation of a method to call at run time

# Virtual methods in C++

---

```
class A { virtual void m(int p, int q) { ... } ... }  
class B : public A { virtual void m(int r, int s) { ... } ... }  
A* x;  
x = new B();  
x->m(1,2);
```

- Since x has declared type A\*, the compile time target is method m in class A
- The run time target is m in B
  - Without the keyword virtual, the run time target will be the same as the compile time target

# Abstract classes

---

- Abstract class: A class that contains abstract methods
  - `abstract void m(int x);` // Java
  - `virtual void m(int x) = 0;` // C++
- Abstract classes allow defining and creating functionality that subclasses can implement or override
- We cannot say `new X()` if `X` is an abstract
  - Why not?
- An abstract method can be the compile time target of a method call
  - But not the run time target
- Sometimes non-abstract classes are referred to as “concrete classes”



# Interfaces


---

- Very similar to abstract classes in which all methods are abstract
  - Functionality can be defined but not implemented
- A reference variable can be of interface type, and can refer to any instance of a class that implements the interface
- An interface method can be the compile time target of a method call
- Java 8 allows default and static methods in interfaces

# Example

```
interface X { void m(); }  
interface Y { void n(); }  
abstract class A implements X {  
    void m() { ... }  
    abstract void m2();  
}  
class B extends A implements Y {  
    void m2() { ... }  
    void n() {...}  
}  
X x = new B(); x.m();  
Y y = new B(); y.n();  
A a = new B(); a.m2();
```

compile-time  
targets



# Outline

---

- Classes and objects
- Methods
  - Inheritance, polymorphism
  - **Static methods and fields**
- Implementation: compilation, allocation
- Types
- Memory and type safety
- Memory Management

# Static methods and fields

---

- **Static field**: A single copy for the entire class
- **Static method**: not invoked on an object
  - Just like a regular subroutine in a procedural language
- **Terminology**:
  - Static method/field = class method/field
  - Instance method/field = non-static method/field

# Static example (Java)

---

```
class X { ...  
    private static int num = 0;  
    // constructor  
    public X() { num++; }  
    public static int numInstances()  
        { return num; }  
}
```

in main:

```
X x1 = new X(); X x2 = new X();  
int n = X.numInstances();
```

 returns 2

# Static example (C++)

---

```
class X { ...  
    private: static int num;  
    public: X();  
    public: static int numInstances();  
}
```

```
int X::num = 0;
```

```
X::X() { num++; }
```

```
int X::numInstances() { return num; }
```

in main:

```
X* x1 = new X; X* x2 = new X;
```

```
int num = X::numInstances();
```

 returns 2

# Example: Singleton pattern (Java)

---

```
class Logger {  
    private Logger() { }  
    private static Logger instance = null;  
    public static Logger getInstance() {  
        if (instance == null)  
            instance = new Logger();  
        return instance;  
    }  
}  
  
client code: Logger.getInstance().writeLog(...)
```

# Outline

---

- Classes and objects
- Methods
  - Inheritance, polymorphism
  - Static methods and fields
- **Implementation: compilation, allocation**
- Types
- Memory and type safety
- Memory Management



# Implementation techniques for Java

---

- The **compiler** takes as input source code
  - Oracle/Sun provides a standard compiler; other can build their own compilers if they want
  - Typically, class A is stored in file A.java
    - Exception: nested classes
- **Compiler output:** Java bytecode
  - A.java -> A.class
  - A standardized platform independent representation of Java code
  - Essentially, a programming language that is understood by the JVM
    - Why not have the JVM directly interpreter Java source?

# Rectangle.class

---

```
class Rectangle extends  
java.lang.Object {  
    public double height;  
    public double width;  
    Rectangle();  
    public double area();  
}
```

```
Rectangle()  
    0 aload_0  
    1 invokespecial #3 <Method  
      java.lang.Object()>  
    4 return  
double area()  
    0 aload_0  
    1 getfield #4 <Field double height>  
    4 aload_0  
    5 getfield #5 <Field double width>  
    8 dmul  
    9 dreturn
```

# Execution model

---

- Java **bytecode** is executed by a Java Virtual Machine (JVM)
  - Oracle/Sun provide several kinds of JVMs for various platforms
    - e.g. Windows, Linux, MacOS, Solaris, ect)
  - Several other vendors also provide JVMs
    - e.g. IBM sells a JVM that is performance-tuned for enterprise server applications
- Platform independence: as long as there are JVMs available, the exact same Java **bytecode** can be executed anywhere

# JVM

---

- There are two ways to execute the **bytecode**
- **Interpretation**: The VM just executes each **bytecode** instruction itself
  - Initial JVMs used this model
- **Compilation**: The VM uses its own internal compiler to translate **bytecode** to native code for the platform
  - The native code is executed by the platform
  - Faster than interpretation

# Compilation inside a VM

---

- **Just-in-time**: The first time some **bytecode** needs to be executed, it is compiled to native code on the fly
  - Typically done at method level, the first time a method is invoked the compiler kicks in
  - Problems: compilation has overhead, and the overall running time may actually increase
- **Profile-driven compilation**
  - Start executing through interpretation, but track “hot spots” (frequently executed methods) and after a certain threshold is reached compile them

# Outline

---

- Classes and objects
- Methods
  - Inheritance, polymorphism
  - Static methods and fields
- Implementation: compilation, allocation
- **Types**
- Memory and type safety
- Memory Management

# Types

---

- Organization of **untyped values**
  - Untyped universes: bit strings, S-expr, ...
  - Categorize based on usage and behavior
- Type = set of computational entities with uniform behavior
- Constraints to enforce correctness
  - Check the applicability of operations
    - Should not try to multiply two strings
    - Should not use a character value as a condition of an if-statement
    - Should not use an integer as a pointer

# Examples of Type Checking

---

- Built-in operators should get operands of correct types
- Type of left-hand side must agree with the value on the right-hand side
- Procedure calls: check number and type of actual arguments
- Return type should match returned value



# Static Typing

---

- Statically typed languages: expressions in the code have **static types**
  - static type = claim about run-time values
  - Types are either **declared** or **inferred**
  - Examples: C, C++, Java, ML, Pascal, Modula-3
- A statically typed language typically does some form of **static type checking**
  - E.g., at compile time Java checks that the [] operator is applied to a value of type “array”
  - May also do dynamic (run-time) checking
    - e.g., Java checks at run time for array indices out of bounds and for null pointers

# Dynamic Typing

---

- Dynamically-typed languages: entities in the code **do not** have static types
  - Examples: Lisp, Scheme, CLOS, Smalltalk, Perl, Python
  - Entities in the code do not have declared types, and the compiler does not try to infer types for them
- **Dynamic type checking**
  - Before an operation is performed at run time
  - E.g., in Scheme: **(+ 5 #t)** fails at run time, when the evaluation expects to see two numeric atoms as operands of +

# Type (and memory) safety

---

- **Type-safe** languages: type-incorrect operations are not performed at run time
  - Things cannot “go wrong”: no undetected type errors
  - Certain run-time errors are possible but clearly marked as such
    - i.e. array index out of bounds, null pointer
  - C/C++: type unsafe
  - Java: type safe
- Independent of static vs. dynamic
  - Lisp, Scheme, Python: dynamically typed; type safe
  - Forth: dynamically typed; type unsafe

# Examples of Types

---

- Integers
- Arrays of integers
- Pointers to integers
- Records with fields **int x** and **int y**
  - e.g., “struct” in C
- Objects of class C or a subclass of C
  - e.g., C++, Java, C#
- Functions from any list to integers

# Numeric Types

---

- Varied from language to language
- C does not specify the ranges of numeric types
  - Integer types: char, short, int, long, long long
    - Includes “unsigned” versions of these
  - Floating-point types: float, double, long double
- Java specifies the ranges of numeric types
  - byte: 8-bit signed two's complement integer [-128,+127]
  - short: 16-bit signed two's complement integer [-32768,+32,767]
  - int: 32-bit signed two's complement integer
    - [-2147483648,+2147483647]
  - long: 64-bit signed two's complement integer
    - [-9223372036854775808, +9223372036854775807]
  - float/double: single/double-precision 32-bit IEEE 754 floating point
  - char: single 16-bit Unicode character; minimum value of '\u0000' (or 0) and a maximum value of '\uffff' (or 65535)

# Enumeration Types

---

- C: a set of named integer constant values
  - Example from the C specification

```
enum hue { chartreuse, burgundy, claret=20, winedark };  
/* the set of integer constant values is { 0, 1, 20, 21 } */  
enum hue col, *cp;  
col = claret; cp = &col;  
if (*cp != burgundy) ...
```
- Java: a fixed set of named items (not integers)

```
enum Day { SUNDAY, MONDAY, TUESDAY, WEDNESDAY, THURSDAY, FRIDAY,  
SATURDAY }
```

  - In reality, it is like a class: e.g., it can contain methods

# Types as Sets of Values

---

- Integers
  - Any number that can be represented in 32 bits in signed two's-complement
  - “**type int**” =  $\{ -2^{31}, \dots, 2^{31} - 1 \}$
- Class type (not the same as a class)
  - Any object of class C or a subclass of C
  - “**type C**” = set of all instances of C or of any transitive subclass of C (“**class C**” is just a blueprint for objects)
- **Subtypes are subsets**: T2 is a **subtype** of T1 if T2's set of values is a subset of T1's set of values

# Monomorphism vs. Polymorphism

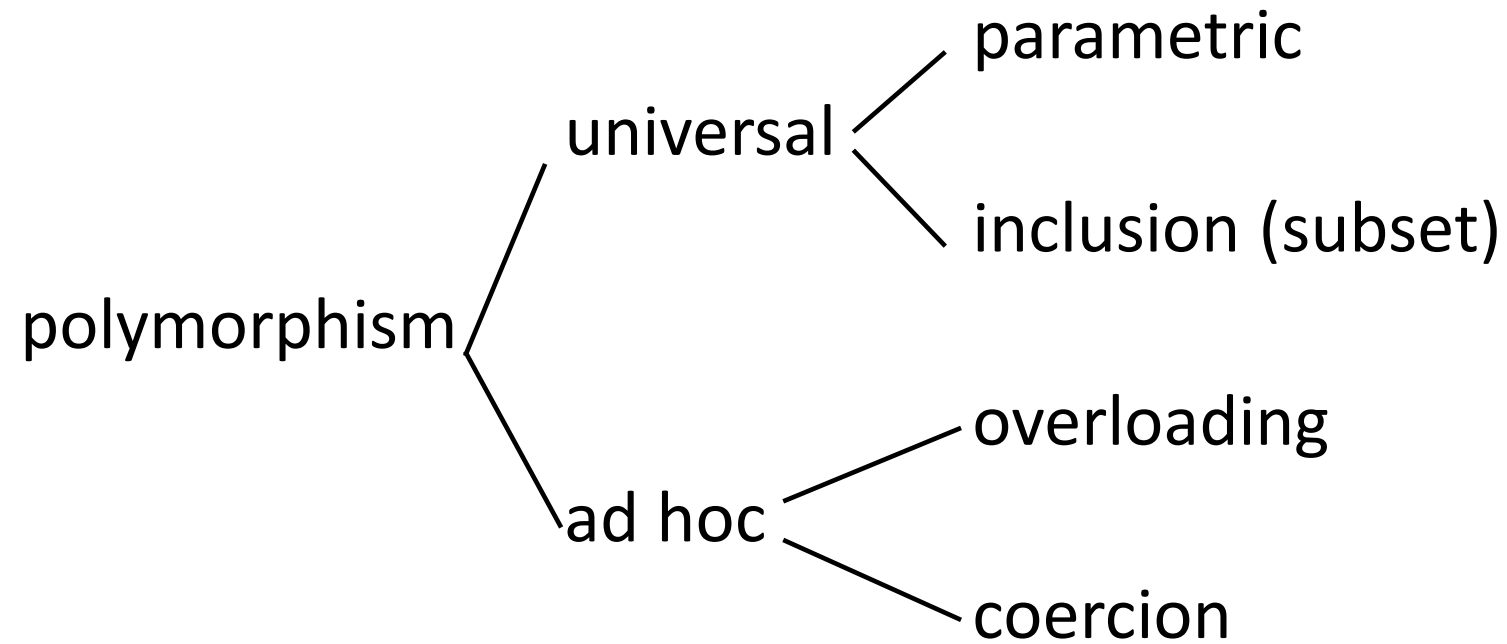
---

- Greek:
  - mono = single
  - poly = many
  - morph = form
- Monomorphism
  - Every computational entity belongs to exactly one type
- Polymorphism
  - A computational entity can belong to multiple types



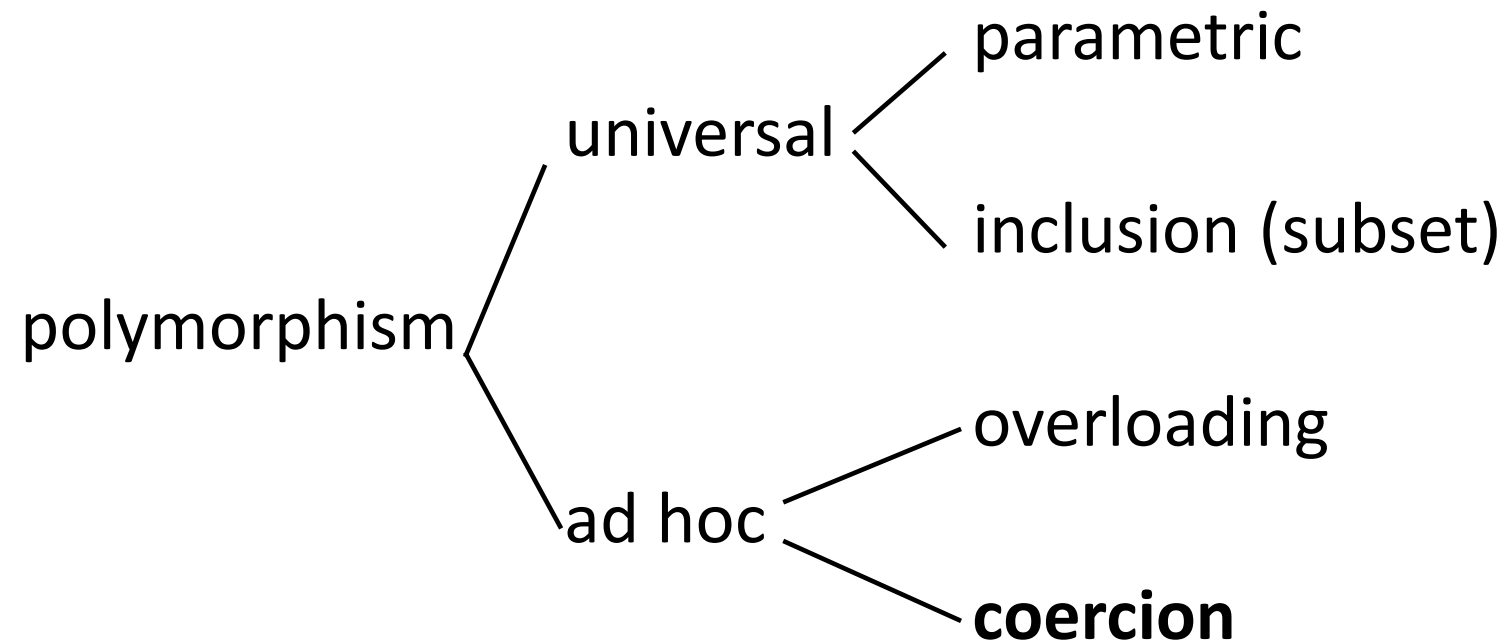
# Types of Polymorphism

---



# Types of Polymorphism

---



# Coercion

---

- Values of one type are silently converted to another type
  - e.g. addition:  $3.0 + 4$  : converts 4 to 4.0
    - $\text{int} \times \text{int} \rightarrow \text{int}$  or  $\text{real} \times \text{real} \rightarrow \text{real}$
- In a context where the type of an expression is not appropriate
  - either an automatic coercion (conversion) to another type is performed automatically
  - or if not possible: compile-time error

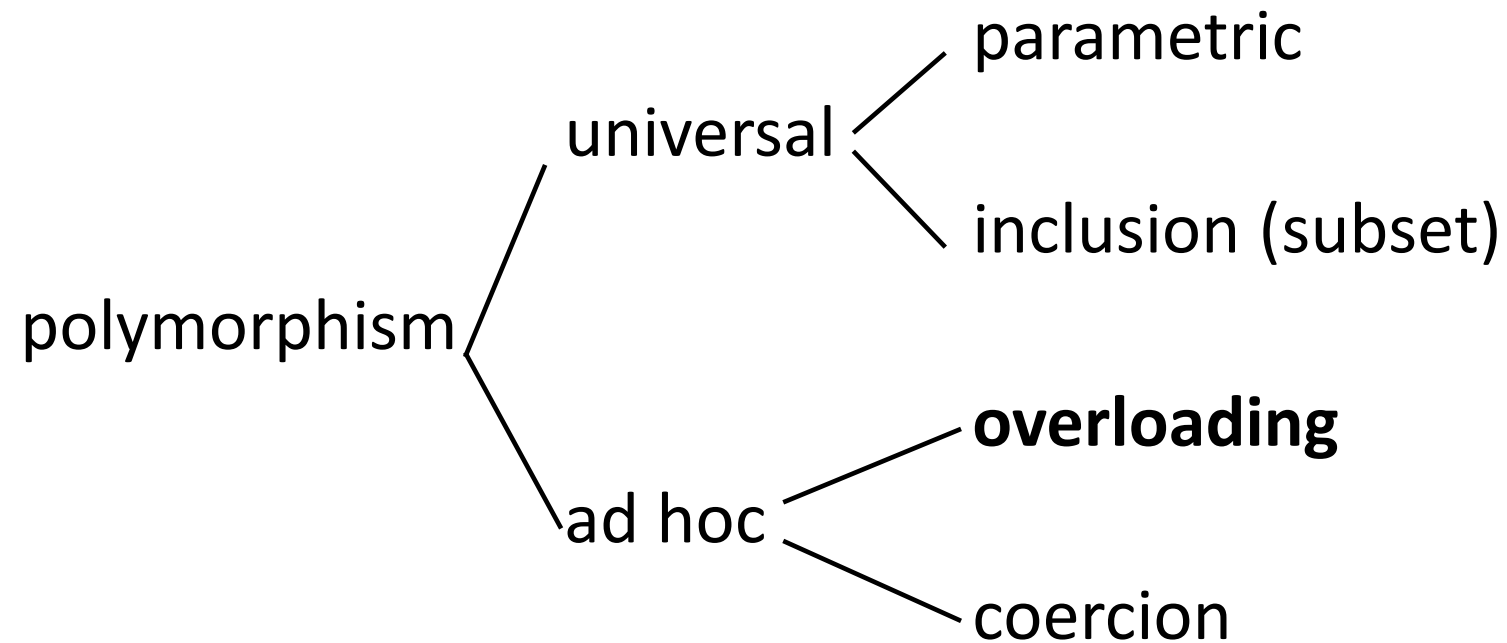
# Coercions

---

- Widening
  - coercing a value into a “larger” type
  - e.g., **int** to **float**, subclass to superclass
- Narrowing
  - coercing a value into a “smaller” type
  - loses information, e.g., **float** to **int**

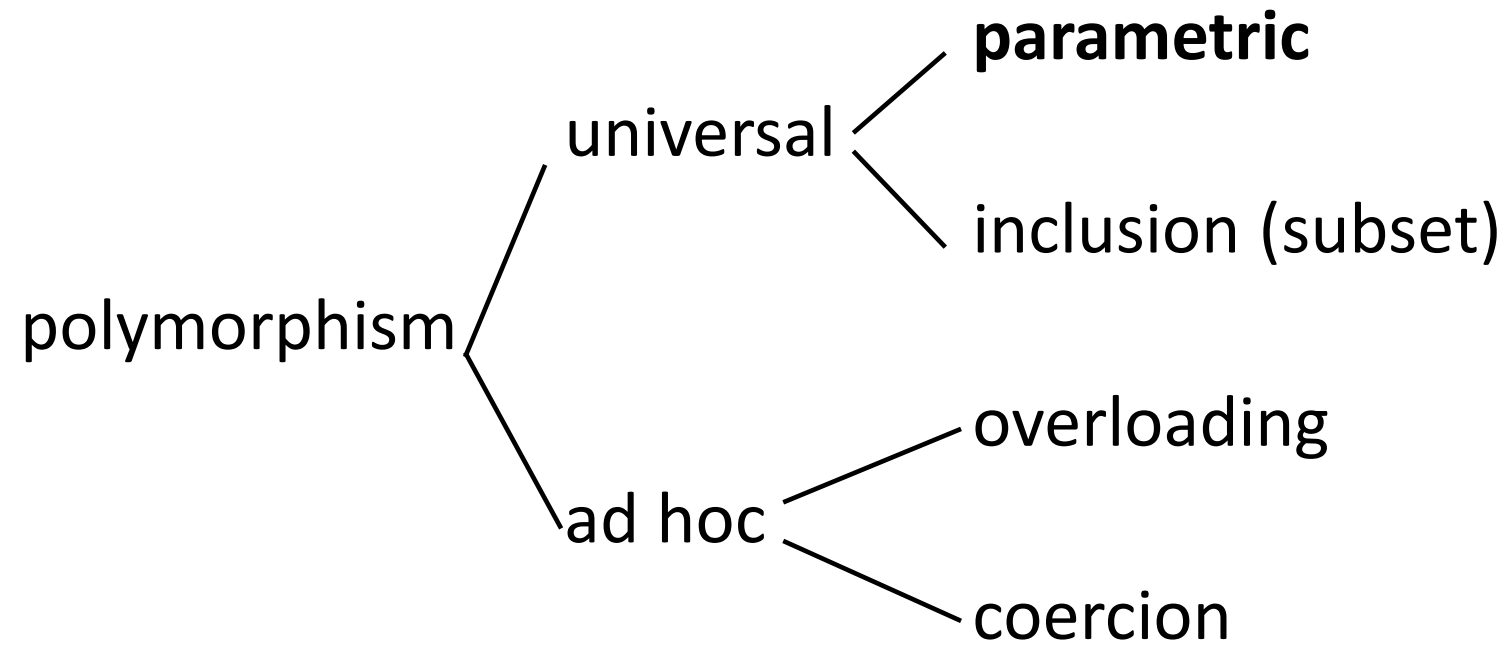
# Types of Polymorphism

---



# Types of Polymorphism

---



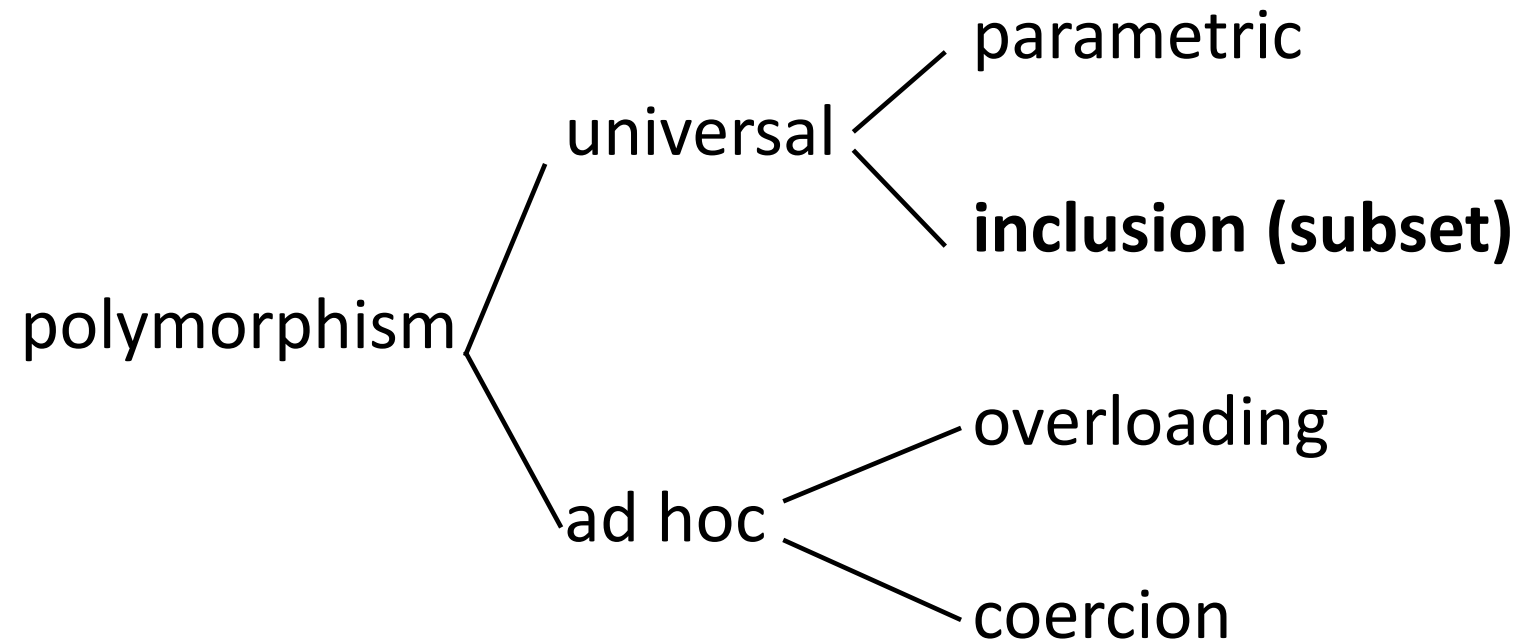
# Parametric Polymorphism: Generics in Java

---

```
package java.util;
public interface Set<E> extends Collection<E> { ...
    Iterator<E> iterator();
    boolean add(E e);
    boolean addAll(Collection<? extends E> c); }
class Rectangle { ... }
class SwissRectangle extends Rectangle { ... }
Set<Rectangle> s = new HashSet<Rectangle>();
s.add(new Rectangle(1.,2.)); s.add(new SwissRectangle(3.,4.,5));
Set<SwissRectangle> s2 = new TreeSet<SwissRectangle>();
s2.add(new SwissRectangle(6.,7.,8)); s.addAll(s2);
```

# Types of Polymorphism

---





# Inclusion (Subset) Polymorphism

---

- Subtype relationships among types
  - Defined by “Y is subset of X” (i.e., set inclusion)
- A computational entity of a subtype may be used in any context that expects an entity of a supertype
- Typical examples
  - Imperative languages: record types
  - Object-oriented languages: class types

# Subtyping in Java

---

- Recall that **class type C** is the set of all instances of class C or of any transitive subclass of C
- Subtyping between class types

```
class X { int m () { ... } }  
class Y extends X { int m () { ... } }  
X x = new Y();  
int i = x.m();
```
- Interface type: the set of all instances of classes that implements the interface (transitively)

```
interface Z { bool m(); }  
class W implements Z { bool m() { ... } }  
Z z = new W(); bool b = z.m();
```

# Outline

---

- Classes and objects
- Methods
  - Inheritance, polymorphism
  - Static methods and fields
- Implementation: compilation, allocation
- Types
- **Memory and type safety**
- Memory Management

# Memory and type safety

---

```
int[] a = new int[64];
```

```
...
```

```
a[82] = ...;
```

# Memory and type safety

---

```
int[] a = new int[64];
```

```
...
```

```
a[82] = ...;
```

C: undefined behavior → major source of security exploits (how?)

Java: throws `ArrayIndexOutOfBoundsException`

How? Instrumentation added by JIT compiler (or proved unnecessary)

# Memory and type safety

---

```
MyType* a = new MyType();
```

```
...
```

```
*((void*)a + 16) = 42;
```

# Memory and type safety

---

```
MyType* a = new MyType();
```

```
...
```

```
*((void*)a + 16) = 42;
```

C++: Undefined behavior

Java: Pointer arithmetic isn't part of the language

# Memory and type safety

---

SomeType\* a = ...;

b = (IncompatibleType\*) a;

b->f = ...;



# Memory and type safety

---

`SomeType* a = ...;`

`b = (IncompatibleType*) a;`

`b->f = ...;`

C: undefined behavior → potential security exploit

Java: throws `ClassCastException`

How? Instrumentation added by JIT compiler (or proves unnecessary)

# Memory and type safety

---

```
MyType* a = new MyType(); /* or: malloc(sizeof(MyType)) */
```

```
...
```

```
delete a; /* or: free(a) */
```

```
...
```

```
a->f = ...;
```

# Memory and type safety

---

```
MyType* a = new MyType();
```

```
...
```

```
delete a;
```

```
...
```

```
a->f = ...;
```

# Memory and type safety

---

```
MyType* a = new MyType();
```

```
...
```

```
delete a;
```

```
...
```

```
a->f = ...;
```

C++: Undefined behavior

Java: Garbage collection → no explicit freeing

# Memory and type safety

---

```
MyType* a = new MyType();
```

```
...
```

```
delete a;
```

```
...
```

```
delete a;
```

# Memory and type safety

---

```
MyType* a = new MyType();
```

```
...
```

```
delete a;
```

```
...
```

```
delete a;
```

C++: Undefined behavior

Java: Garbage collection → no explicit freeing

# Memory and type safety

---

```
MyType* a = new MyType();
```

```
container->data = a;
```

```
...
```

```
container->data = NULL; /* Last ptr to a is lost */
```

# Memory and type safety

---

```
MyType* a = new MyType();  
container->data = a;  
...  
container->data = NULL; /* Last ptr to a is lost */
```

C++: Memory leak

Java: Garbage collection

How? Knows all reference types and all “roots”; approximates liveness via transitive reachability



# Outline

---

- Classes and objects
- Methods
  - Inheritance, polymorphism
  - Static methods and fields
- Implementation: compilation, allocation
- Types
- Memory and type safety
- **Memory Management**

# Outline for memory management

---

- **Concepts:**
  - Live vs. dead objects
  - Explicit (manual) MM vs. automatic MM (GC)
  - GC overapproximates live set
  - Memory leaks
- Overview of GC, esp. tracing GC
- GC algorithms and strategies:
  - Mark-sweep GC
  - Copying GC
  - Generational GC

# Live and dead objects

---

- **Live** object will be used in the future
- Other objects are **dead**
- Deallocate as soon as possible after last use (but not before!)

Memory management: Deallocate the **dead** objects in a “timely” fashion

# Explicit (manual) memory management

---

- More code to maintain
- Requires global reasoning
- Correctness
  - Free a live object
  - Free a dead object too late (or never)
- Efficiency can be very high
  - Gives programmers more control over the run-time behavior of the program

# Automatic memory management (garbage collection)

---

- Integral for memory and type safety
  - Protects against some classes of memory errors (dangling pointers, double frees)
  - Essential for Java, C#, PHP, JavaScript, ...
- Reduces programmer burden
- Not perfect, memory can still leak
  - Programmers still need to eliminate all pointers to objects the program no longer needs
- (A mostly solved) challenge: performance

# What is Garbage?

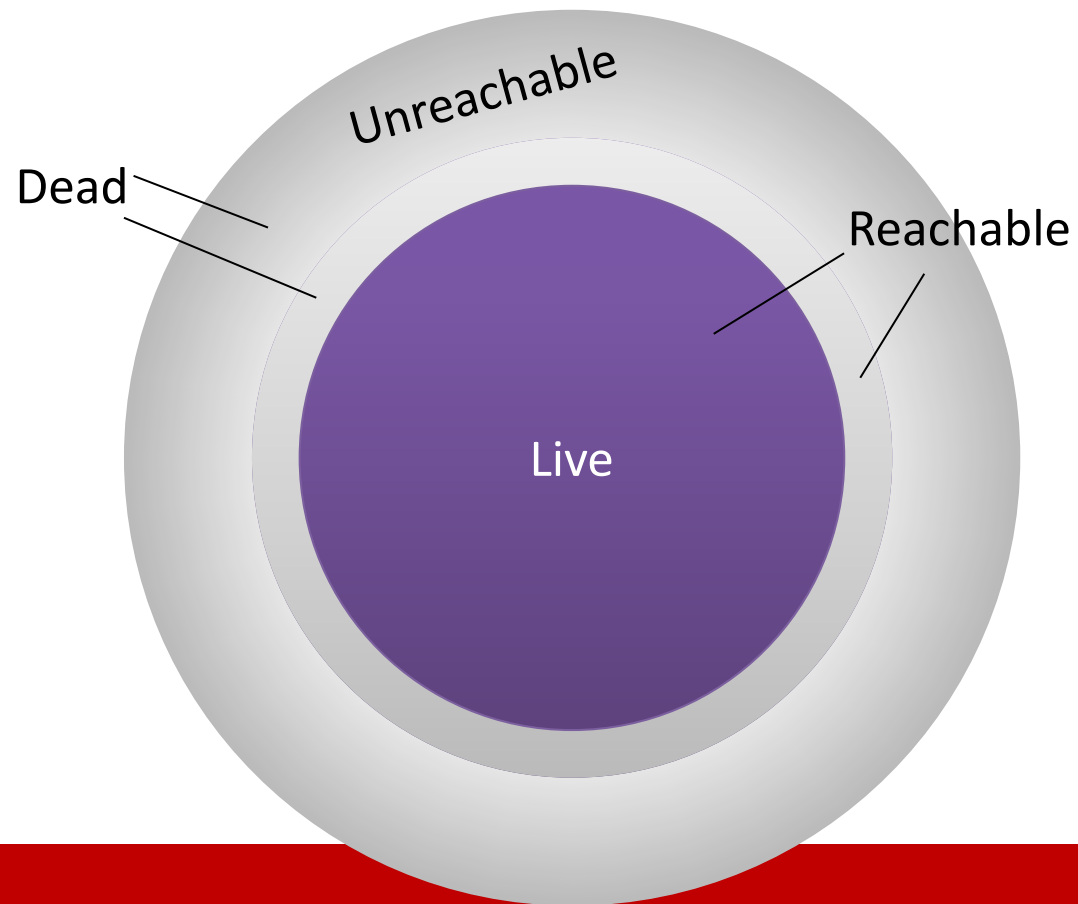
---

- In theory, any object the program will never reference again
  - But compiler & runtime system cannot figure that out
- In practice, any object the program cannot reach is garbage
  - Approximate **liveness** with **reachability**
- Managed languages couple GC with “safe” pointers
  - Programs may not access arbitrary addresses in memory (e.g., Java/C# vs. C/C++)
  - The compiler can identify and provide to the garbage collector all the pointers, thus enforcing “Once garbage, always garbage”
  - Runtime system can move objects by updating pointers

# Liveness under-approximates reachability


(Reachability over-approximates liveness)

---



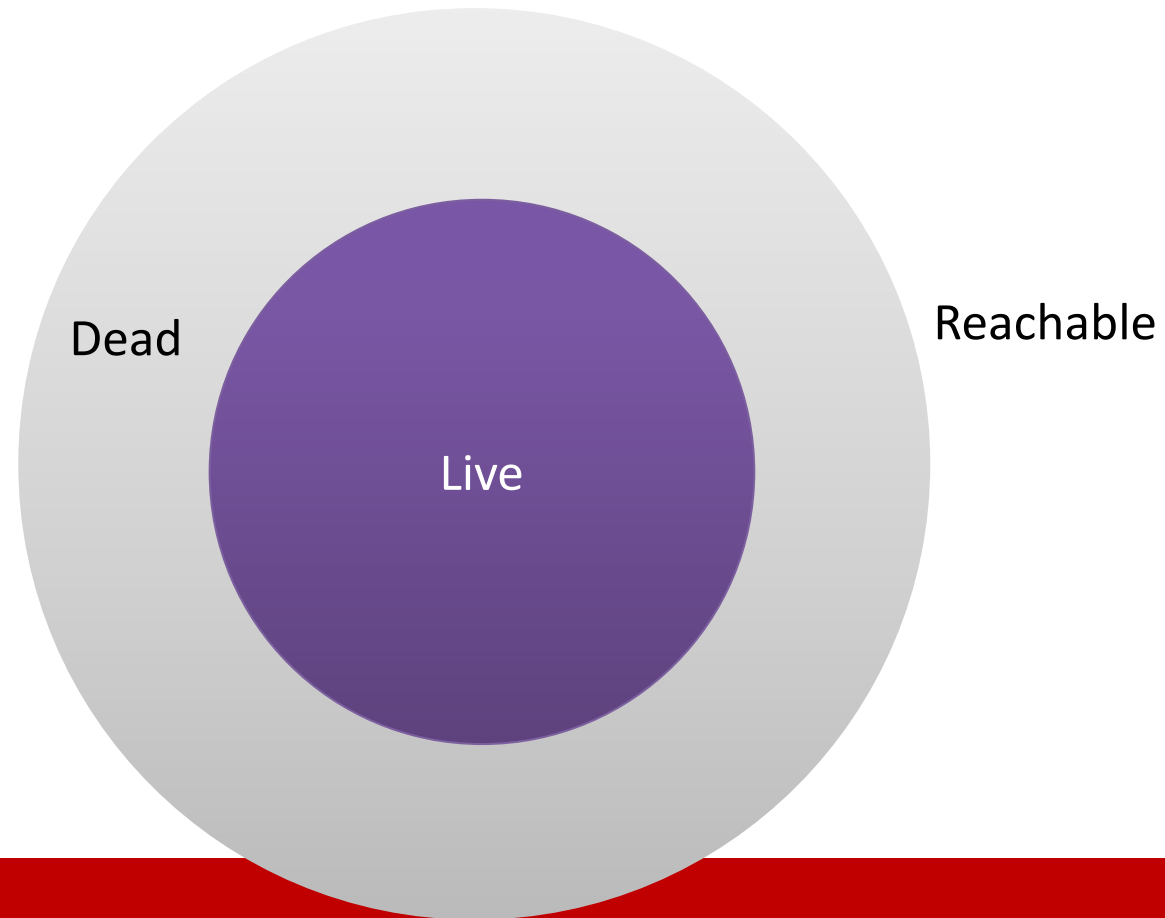
Liveness under-approximates reachability  
(Reachability over-approximates liveness)

---

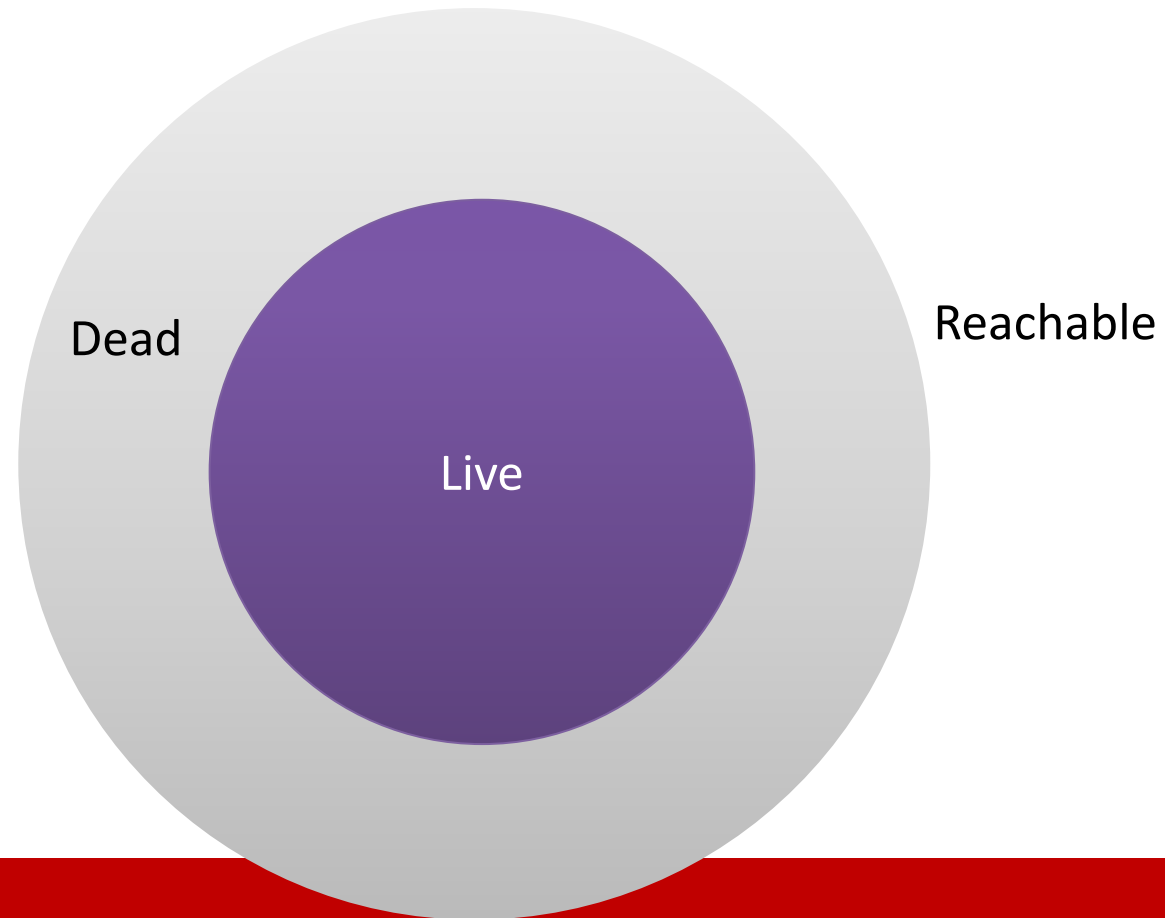




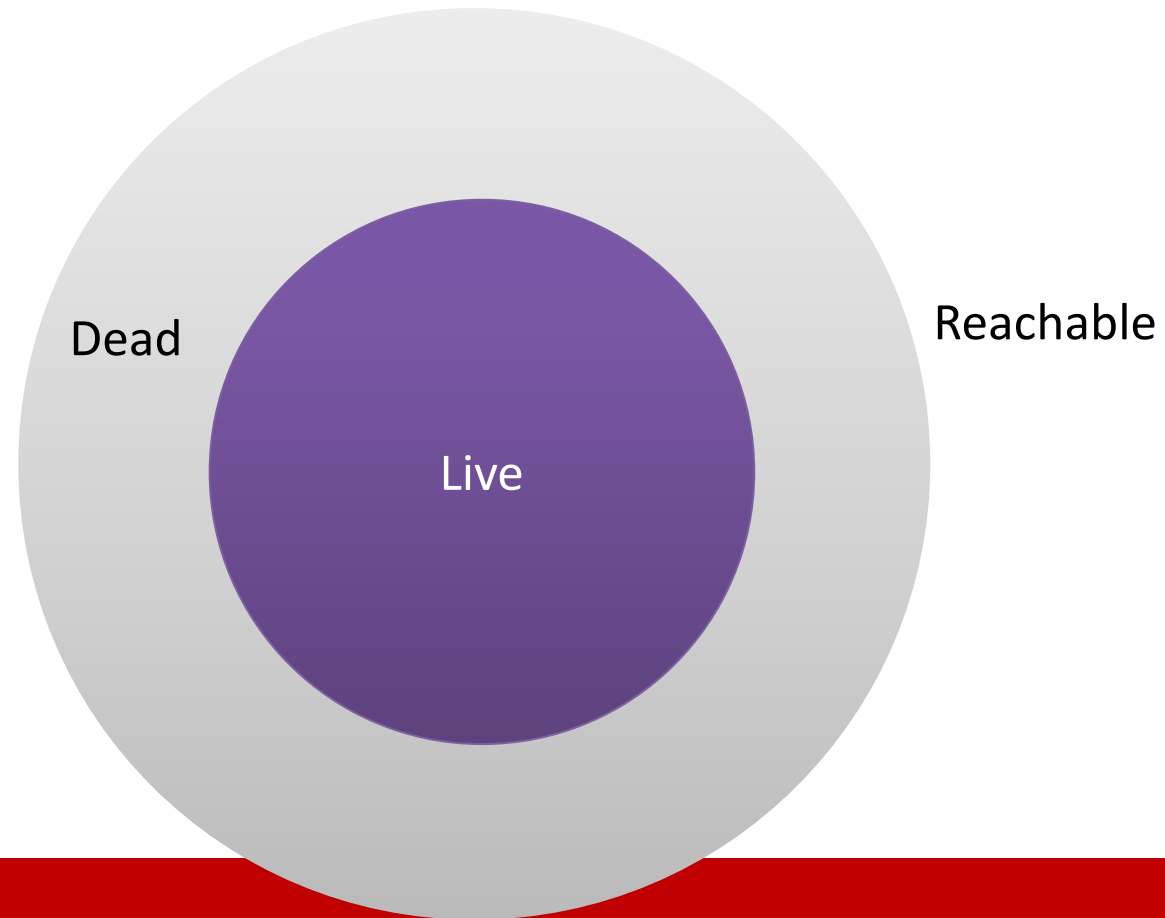
- Can leaks happen in GC'd languages?




- Can leaks happen in GC'd languages?




- Can leaks happen in GC'd languages?



- Can leaks happen in GC'd languages?



- Can leaks happen in GC'd languages?



- Can leaks happen in GC'd languages?

Dead



Live

# Memory leak example

<http://www.codeproject.com/KB/showcase/IfOnlyWedUsedANTSPProfiler.aspx>

---

- Driverless truck
  - 10,000 lines of C#
- Leak: past obstacles remained reachable



# Memory leak example

<http://www.codeproject.com/KB/showcase/IfOnlyWedUsedANTSPProfiler.aspx>

---

- Driverless truck
  - 10,000 lines of C#
- Leak: past obstacles remained reachable
- No immediate symptoms

“This problem was pernicious because it only showed up after 40 minutes to an hour of driving around and collecting obstacles.”





# Memory leak example

<http://www.codeproject.com/KB/showcase/IfOnlyWedUsedANTSPProfiler.aspx>

---

- Driverless truck
  - 10,000 lines of C#
- Leak: past obstacles remained reachable
- No immediate symptoms

“This problem was pernicious because it only showed up after 40 minutes to an hour of driving around and collecting obstacles.”
- Quick “fix”: restart after 40 minutes



# Memory leak example

<http://www.codeproject.com/KB/showcase/IfOnlyWedUsedANTSPProfiler.aspx>

---

- Driverless truck
  - 10,000 lines of C#
- Leak: past obstacles remained reachable
- No immediate symptoms

“This problem was pernicious because it only showed up after 40 minutes to an hour of driving around and collecting obstacles.”
- Quick “fix”: restart after 40 minutes
- Environment sensitive
  - More obstacles in deployment
  - Unresponsive after 28 minutes



# Outline for memory management

---

- Concepts:
  - Live vs. dead objects
  - Explicit (manual) MM vs. automatic MM (GC)
  - GC overapproximates live set
  - Memory leaks
- **Overview of GC, esp. tracing GC**
- GC algorithms and strategies:
  - Mark-sweep GC
  - Copying GC
  - Generational GC

# GC basics

---

Two types (duals of each other):

- Reference counting
  - Work proportional to dead objects
  - Memory freed immediately
  - Cycles are problematic
- Tracing
  - Work proportional to live objects
  - Freeing postponed
  - Can be concurrent

# How does tracing GC work?

---

## Roots

- Local variables: registers & stack locations
- Static variables

## Transitive closure

Memory & type safety ensure GC knows the roots and references exactly

# How does tracing GC work?

---

When does it happen?


- Stop-the-world: safe points inserted by VM
- Concurrent
- Incremental

How many GC threads?

- Single-threaded
- Parallel


# Reachability

- The runtime memory management system examines all global variables, stack variables, and live registers that could refer to heap objects (i.e., the **roots** of reachability)
- GC threads can **trace** these pointers through the heap (following object fields that themselves point to heap objects) to find all reachable objects



# Reachability


- Tracing collector
  - Marks the objects reachable from the roots as **live objects**, and then performs a **reachability** computation from them






# Reachability

- Tracing collector
  - Marks the objects reachable from the roots as **live objects**, and then performs a **reachability** computation from them




# Reachability

- Tracing collector
  - Marks the objects reachable from the roots as **live objects**, and then performs a **reachability** computation from them




# Reachability

- Tracing collector
  - Marks the objects reachable from the roots as **live objects**, and then performs a **reachability** computation from them
- All unmarked objects are **dead**



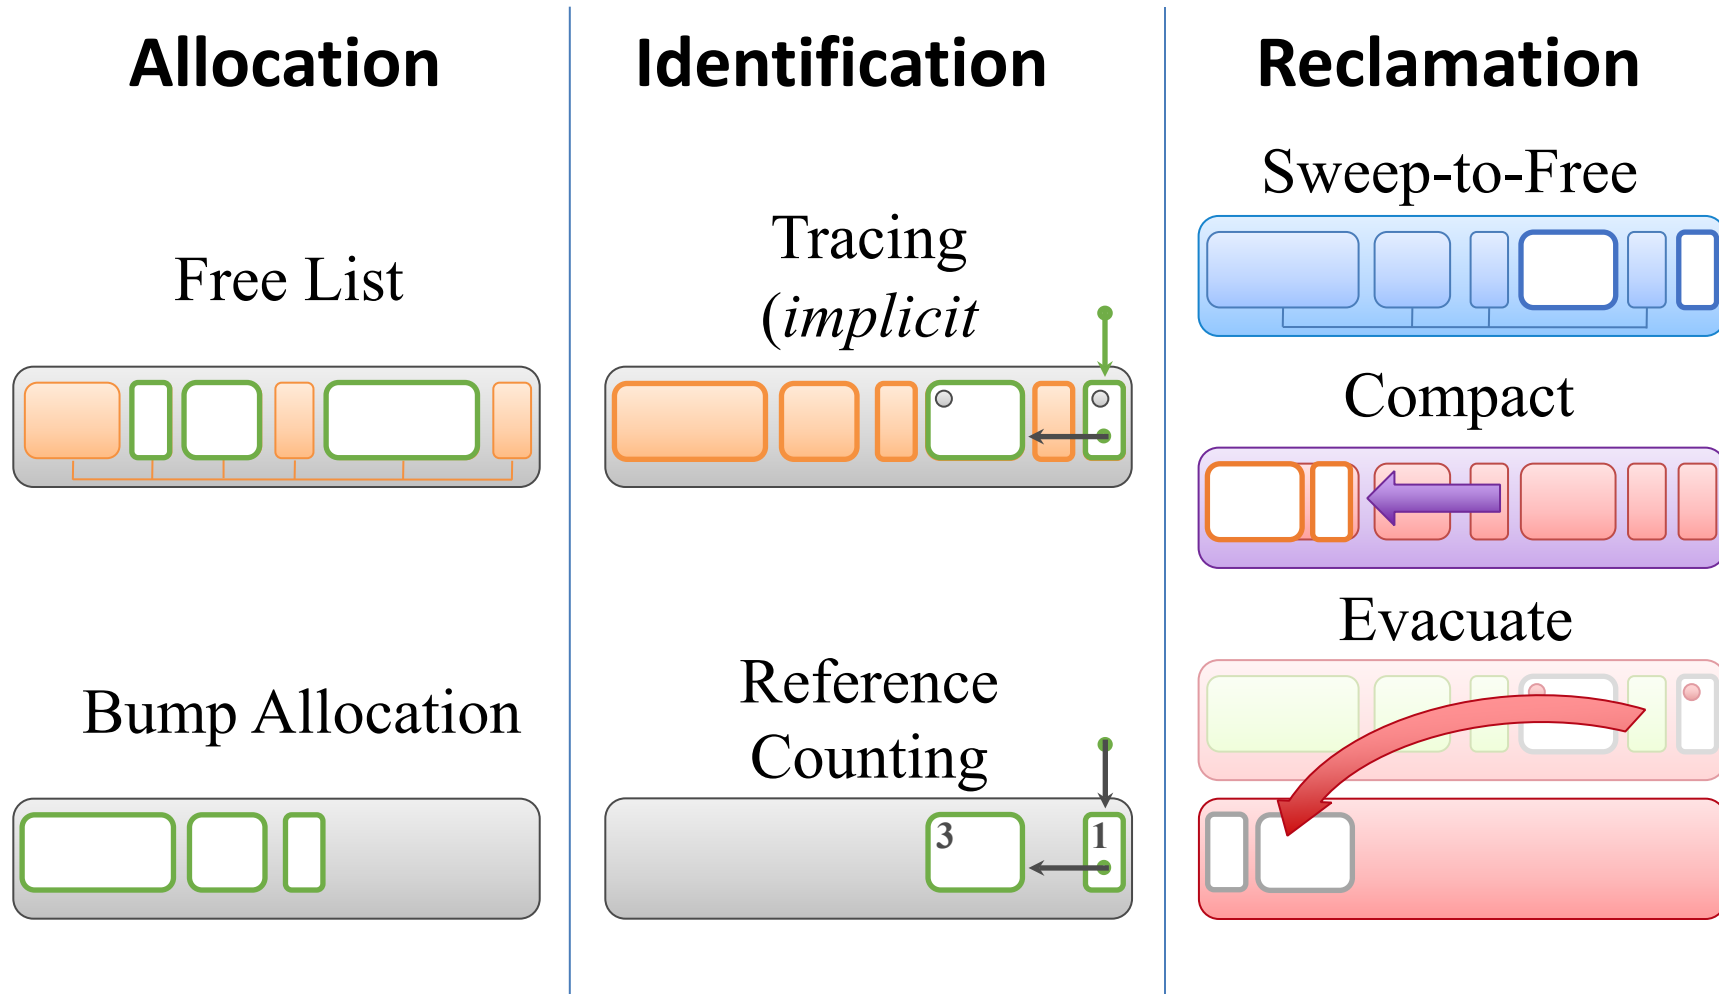
# Reachability

- Tracing collector
  - Marks the objects reachable from the roots as **live objects**, and then performs a **reachability** computation from them
- All unmarked objects are **dead**



# The Big Picture

- Heap organization; basic algorithmic components

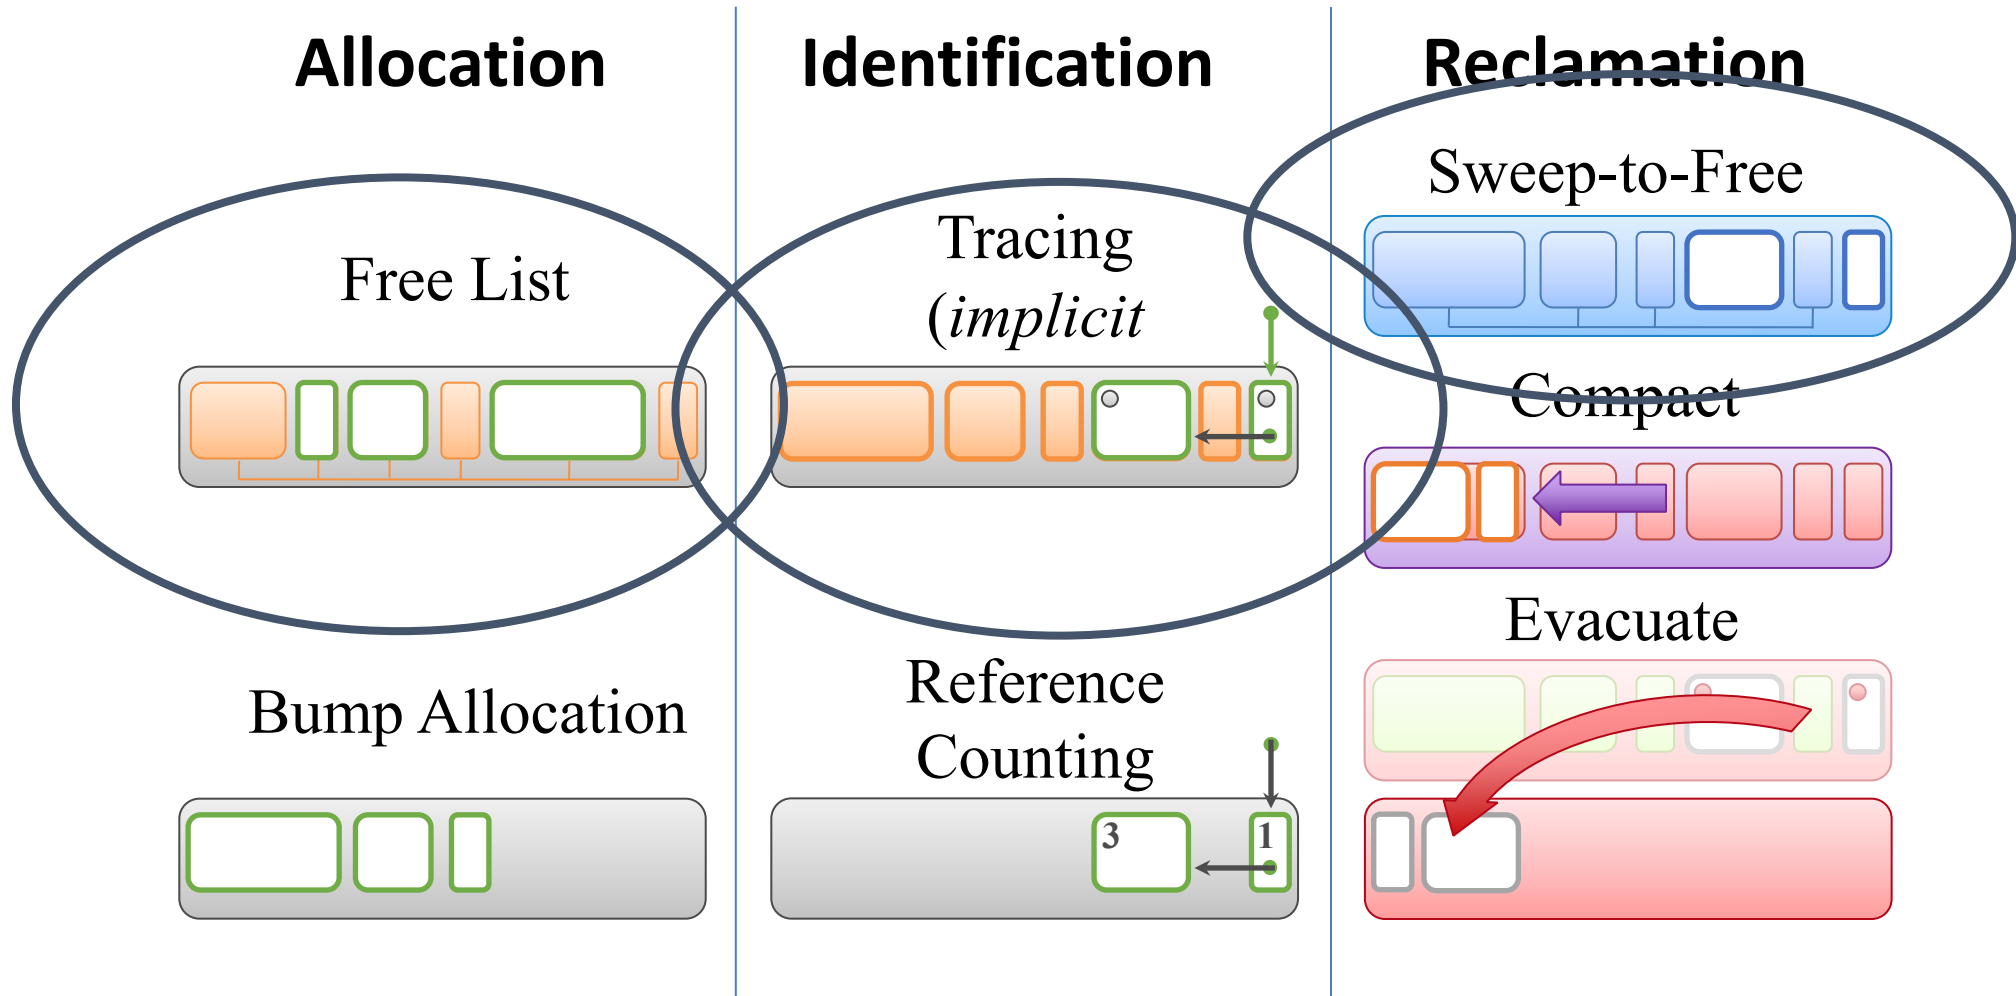


# Outline for memory management

- Concepts:
  - Live vs. dead objects
  - Explicit (manual) MM vs. automatic MM (GC)
  - GC overapproximates live set
  - Memory leaks
- Overview of GC, esp. tracing GC
- GC algorithms and strategies:
  - **Mark-sweep GC**
  - Copying GC
  - Generational GC


# The Big Picture

- Heap organization; basic algorithmic components



# Mark-and-Sweep Implementation


- Free-lists organized by size
  - blocks of same size, or
  - individual objects of same size
- Most objects are small  $< 128$  bytes






# Mark-and-Sweep Implementation

- Allocation
  - Grab a free object off the free list




# Mark-and-Sweep Implementation

- Allocation
  - Grab a free object off the free list





# Mark-and-Sweep Implementation

- Allocation
  - Grab a free object off the free list





# Mark-and-Sweep Implementation

- Allocation
  - Grab a free object off the free list
  - If there is no more memory of the right size, a garbage collection is triggered
  - Mark phase - find the live objects
  - Sweep phase - put free ones on the free list





# Mark-and-Sweep Implementation

- Mark phase
  - Reachability computation on the heap, marking all live objects
- Sweep phase
  - Sweep the memory for free objects, and populate the free lists





# Mark-and-Sweep Implementation

- Mark phase
  - Reachability computation on the heap, marking all live objects
- Sweep phase
  - Sweep the memory for free objects, and populate the free lists




# Mark-and-Sweep Implementation

- Mark phase
  - Reachability computation on the heap, marking all live objects
- Sweep phase
  - Sweep the memory for free objects, and populate the free lists




# Mark-and-Sweep Implementation

- Mark phase
  - Reachability computation on the heap, marking all live objects
- Sweep phase
  - Sweep the memory for free objects, and populate the free lists



free lists



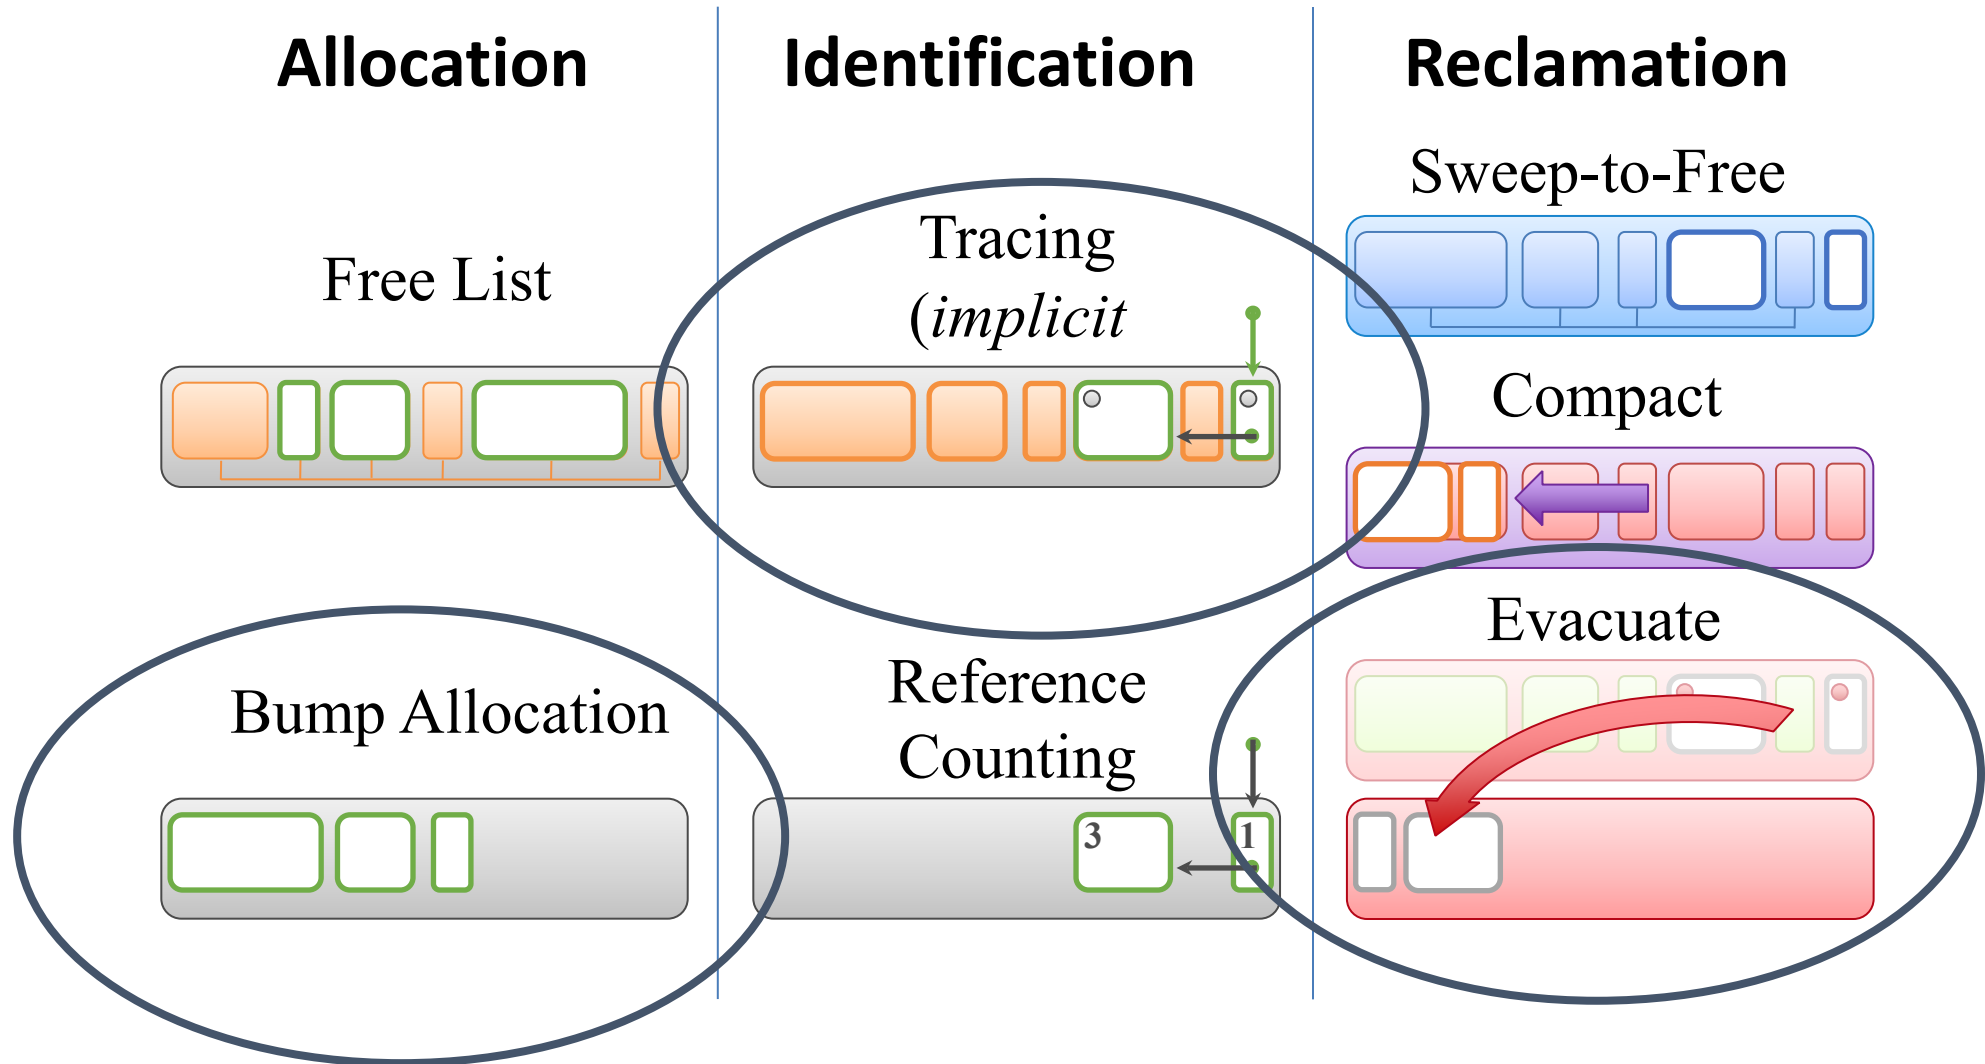


# Outline for memory management

- Concepts:
  - Live vs. dead objects
  - Explicit (manual) MM vs. automatic MM (GC)
  - GC overapproximates live set
  - Memory leaks
- Overview of GC, esp. tracing GC
- GC algorithms and strategies:
  - Mark-sweep GC
  - **Copying GC**
  - Generational GC


# The Big Picture

- Heap organization; basic algorithmic components




# Semispace

- Fast **bump pointer** allocation
- Requires copying collection
- Cannot incrementally reclaim memory, must free en masse
- Reserves 1/2 the heap to copy in to, in case all objects are live




# Semispace

- Fast **bump pointer** allocation
- Requires copying collection
- Cannot incrementally reclaim memory, must free en masse
- Reserves 1/2 the heap to copy in to, in case all objects are live




# Semispace

- Fast **bump pointer** allocation
- Requires copying collection
- Cannot incrementally reclaim memory, must free en masse
- Reserves 1/2 the heap to copy in to, in case all objects are live




# Semispace

- Fast **bump pointer** allocation
- Requires copying collection
- Cannot incrementally reclaim memory, must free en masse
- Reserves 1/2 the heap to copy in to, in case all objects are live




# Semispace

- Mark phase:
  - copies object when collector first encounters it
  - installs **forwarding pointers**



# Semispace


- Mark phase:
  - copies object when collector first encounters it
  - installs **forwarding pointers**
  - performs transitive closure, updating pointers as it goes






# Semispace

- Mark phase:
  - copies object when collector first encounters it
  - installs **forwarding pointers**
  - performs transitive closure, updating pointers as it goes




# Semispace

- Mark phase:
  - copies object when collector first encounters it
  - installs **forwarding pointers**
  - performs transitive closure, updating pointers as it goes




# Semispace

- Mark phase:
  - copies object when collector first encounters it
  - installs **forwarding pointers**
  - performs transitive closure, updating pointers as it goes
  - reclaims “from space” en masse




# Semispace

- Mark phase:
  - copies object when collector first encounters it
  - installs **forwarding pointers**
  - performs transitive closure, updating pointers as it goes
  - reclaims “from space” en masse
  - start allocating again into “to space”




# Semispace

- Mark phase:
  - copies object when collector first encounters it
  - installs **forwarding pointers**
  - performs transitive closure, updating pointers as it goes
  - reclaims “from space” en masse
  - start allocating again into “to space”



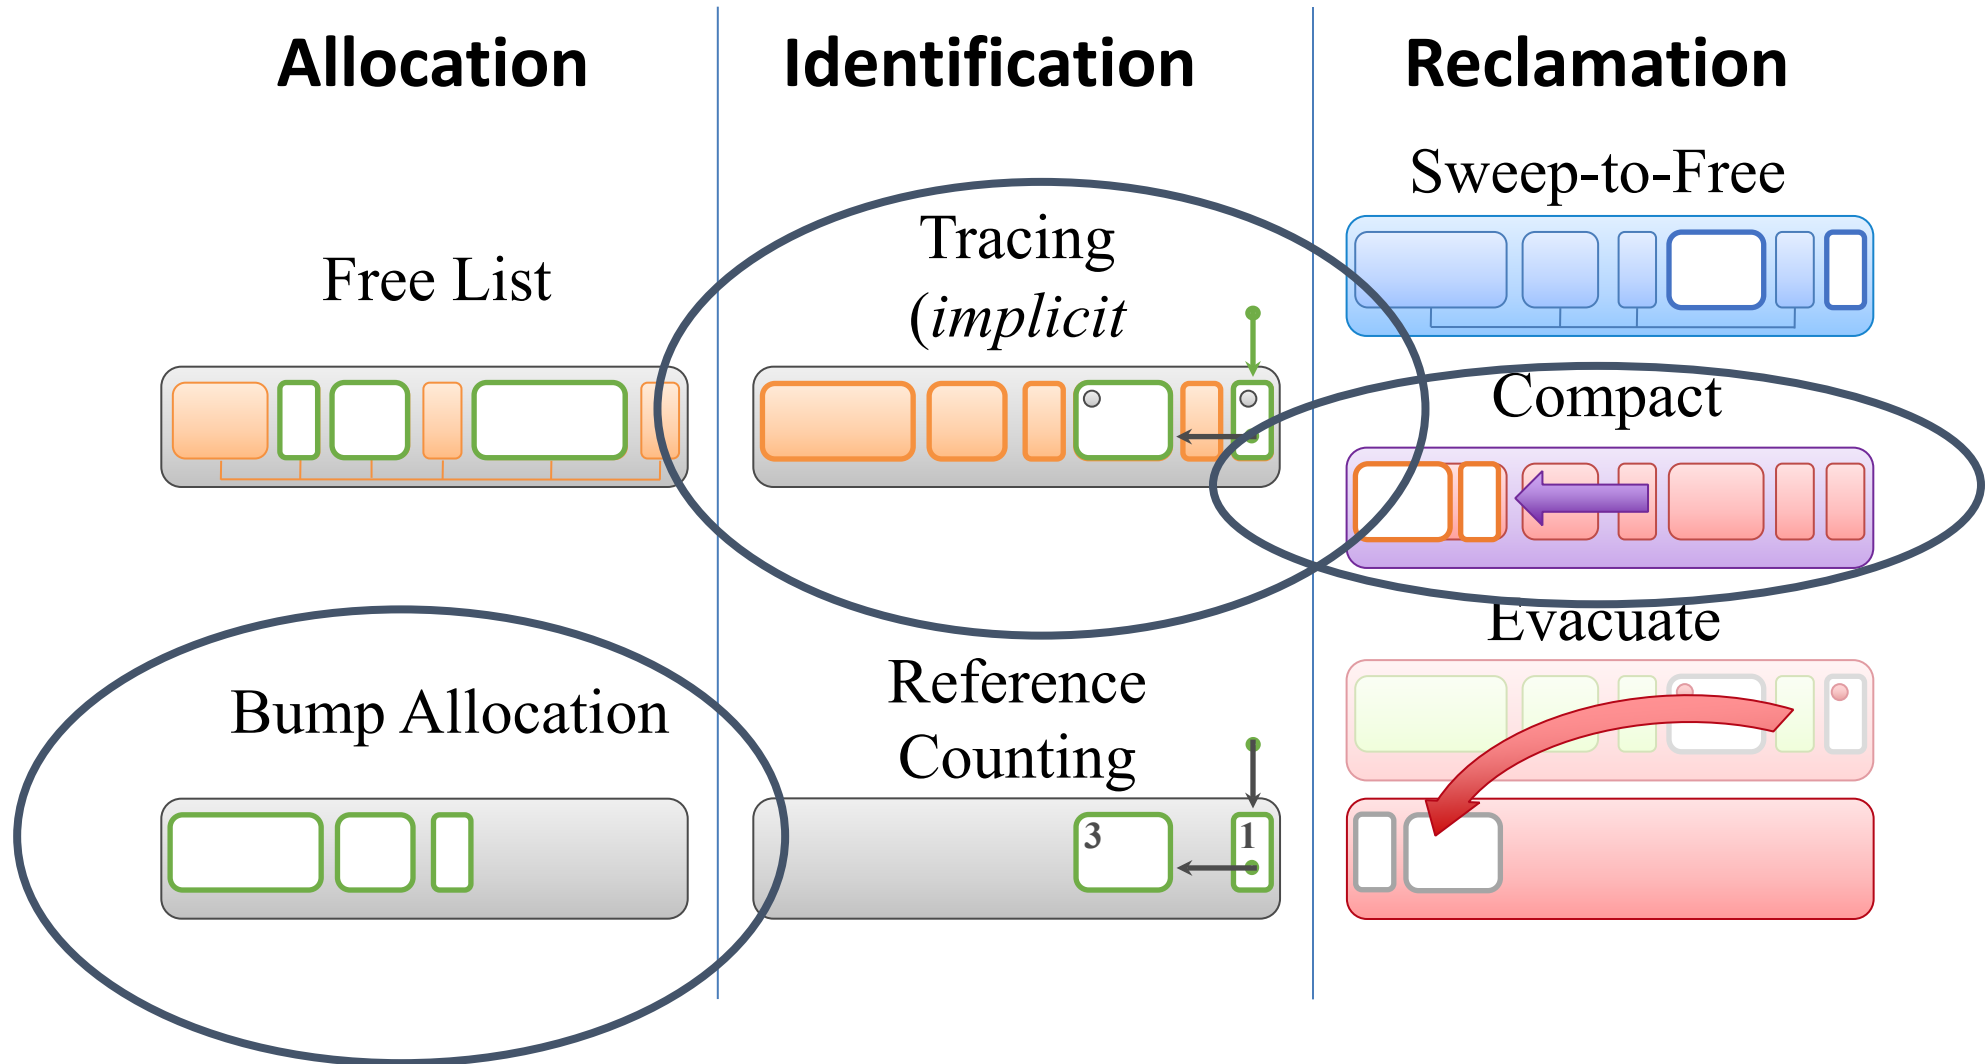
# Semispace

- Notice:
  - fast allocation
  - locality of contemporaneously allocated objects
  - locality of objects connected by pointers
  - wasted space



# The Big Picture

- Heap organization; basic algorithmic components



# Generational Collection

What objects should we put where?

- **Generational hypothesis**
  - young objects die more quickly than older ones [Lieberman & Hewitt'83, Ungar'84]
  - most pointers are from younger to older objects [Appel'89, Zorn'90]
- Organize the heap in to young and old, collect young objects preferentially



# Generational Heap Organization

- **Divide the heap in to two spaces: young and old**
- Allocate in to the young space
- When the young space fills up,
  - collect it, copying into the old space
- When the old space fills up
  - collect both spaces
  - Generalizing to m generations
    - if space  $n < m$  fills up, collect  $n$  through 1




Young



Old

# Generational Heap Organization

- Divide the heap in to two spaces: young and old
- **Allocate in to the young space**
- When the young space fills up,
  - collect it, copying into the old space
- When the old space fills up
  - collect both spaces
  - Generalizing to m generations
    - if space  $n < m$  fills up, collect  $n$  through 1




Young



Old

# Generational Heap Organization

- Divide the heap in to two spaces: young and old
- Allocate in to the young space
- **When the young space fills up,**
  - **collect it, copying into the old space**
- When the old space fills up
  - collect both spaces
  - Generalizing to m generations
    - if space  $n < m$  fills up, collect  $n$  through 1




Young




Old

# Generational Heap Organization

- Divide the heap in to two spaces: young and old
- Allocate in to the young space
- **When the young space fills up,**
  - **collect it, copying into the old space**
- When the old space fills up
  - collect both spaces
  - Generalizing to m generations
    - if space  $n < m$  fills up, collect  $n$  through 1



Young



Old

# Generational Heap Organization

- Divide the heap in to two spaces: young and old
- Allocate in to the young space
- **When the young space fills up,**
  - **collect it, copying into the old space**
- When the old space fills up
  - collect both spaces
  - Generalizing to m generations
    - if space  $n < m$  fills up, collect  $n$  through 1



Young



Old

# Generational Heap Organization

- Divide the heap in to two spaces: young and old
- **Allocate in to the young space**
- When the young space fills up,
  - collect it, copying into the old space
- When the old space fills up
  - collect both spaces
  - Generalizing to m generations
    - if space  $n < m$  fills up, collect  $n$  through 1




Young



Old

# Generational Heap Organization

- Divide the heap in to two spaces: young and old
- Allocate in to the young space
- **When the young space fills up,**
  - **collect it, copying into the old space**
- When the old space fills up
  - collect both spaces
  - Generalizing to m generations
    - if space  $n < m$  fills up, collect  $n$  through 1




Young



Old

# Generational Heap Organization

- Divide the heap in to two spaces: young and old
- Allocate in to the young space
- **When the young space fills up,**
  - **collect it, copying into the old space**
- When the old space fills up
  - collect both spaces
  - Generalizing to m generations
    - if space  $n < m$  fills up, collect  $n$  through 1



Young




Old



# Generational Heap Organization

- Divide the heap in to two spaces: young and old
- Allocate in to the young space
- **When the young space fills up,**
  - **collect it, copying into the old space**
- When the old space fills up
  - collect both spaces
  - Generalizing to m generations
    - if space  $n < m$  fills up, collect  $n$  through 1



Young



Old

# Generational Heap Organization

- Divide the heap in to two spaces: young and old
- Allocate in to the young space
- **When the young space fills up,**
  - **collect it, copying into the old space**
- When the old space fills up
  - collect both spaces
  - Generalizing to m generations
    - if space  $n < m$  fills up, collect  $n$  through 1




Young




Old

# Generational Heap Organization

- Divide the heap in to two spaces: young and old
- Allocate in to the young space
- When the young space fills up,
  - collect it, copying into the old space
- **When the old space fills up**
  - **collect both spaces**
  - Generalizing to m generations
    - if space  $n < m$  fills up, collect  $n$  through 1



Young




Old

# Generational Write Barrier

Unidirectional barrier


- **record only older to younger pointers**
- **no need to record younger to older pointers, since we never collect the old space independently**
- most pointers are from younger to older objects [Appel'89, Zorn'90]
- track the **barrier** between young objects and old spaces



# Generational Write Barrier

unidirectional boundary barrier

```
// original program    // compiler support for incremental collection
p.f = o;               if (p > barrier && o < barrier) {
                        remsetnursery = remsetnursery U &p.f;
                        }
                        p.f = o;
```



# Generational Write Barrier

## Unidirectional

- **record only older to younger pointers**
- **no need to record younger to older pointers, since we never collect the old space independently**
- most pointers are from younger to older objects [Appel'89, Zorn'90]
- most mutations are to young objects [Stefanovic et al.'99]

